

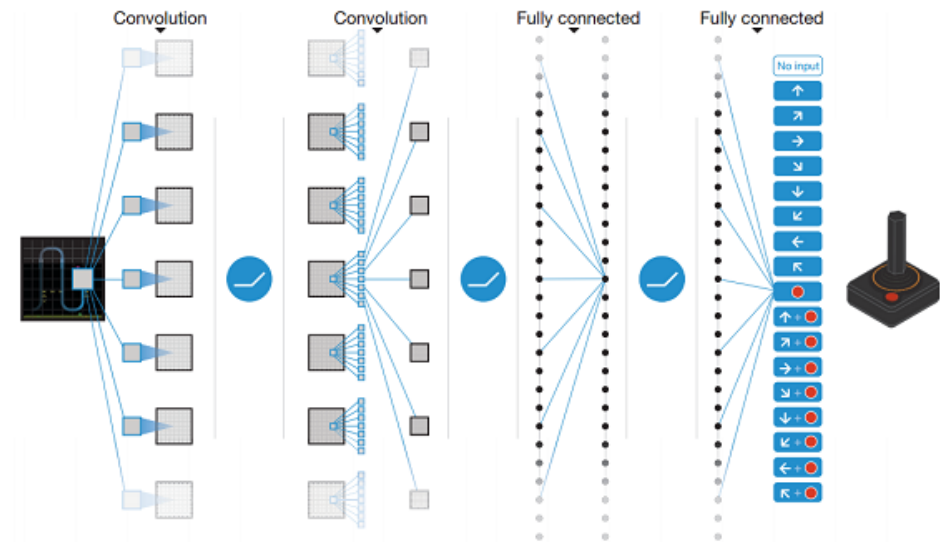
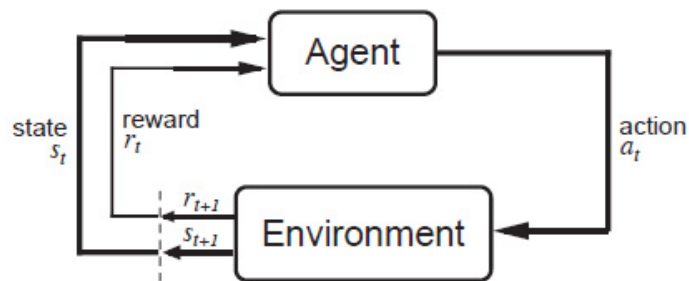
DSCI 552, Machine Learning for Data Science

University of Southern California

M. R. Rajati, PhD

Lesson 13

Reinforcement Learning



Overview

- Supervised Learning: Immediate feedback (labels provided for every input).
- Unsupervised Learning: No feedback (no labels provided).
- Reinforcement Learning: Delayed scalar feedback (a number called reward).

Overview


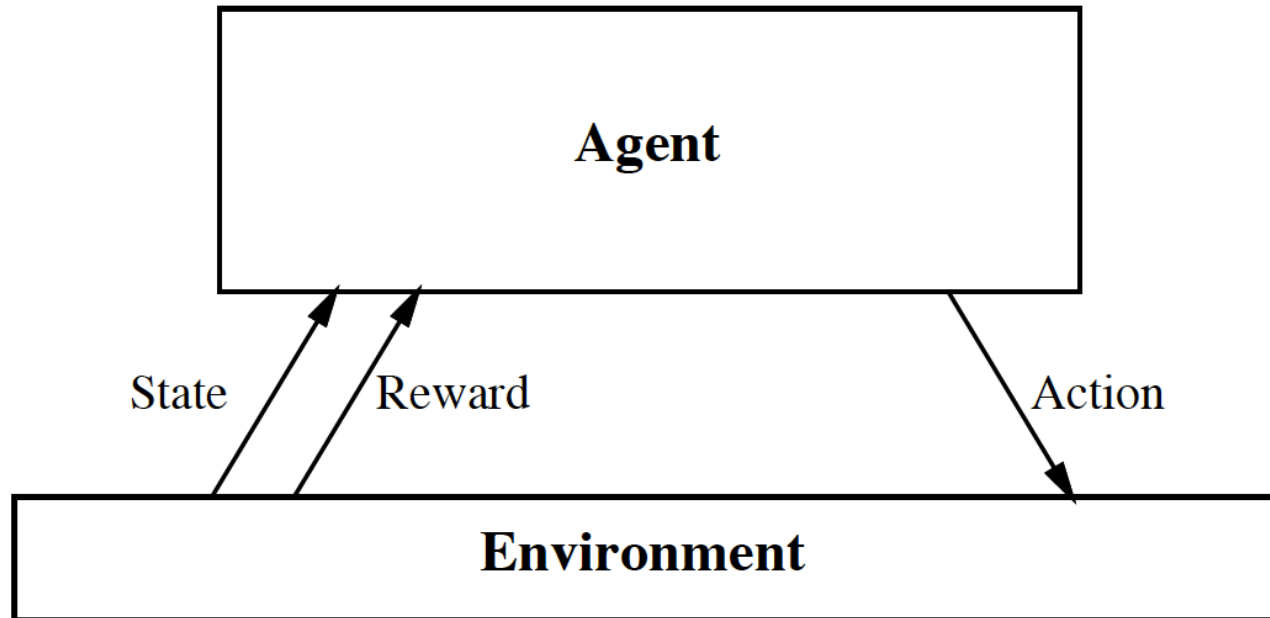
- RL deals with agents that must sense & act upon their environment.
- This combines classical agent-based AI and machine learning techniques.

It is a very comprehensive problem setting.

Overview

- Examples:
 - A robot cleaning my room and recharging its battery
 - Robot-soccer
 - How to invest in shares
 - Modeling the economy through rational agents
 - Learning how to fly a helicopter
 - Scheduling planes to their destinations
 - and so on

The Big Picture



Your action influences the state of the world which determines its reward

Complications

- The outcome of your actions may be uncertain
- You may not be able to perfectly sense the state of the world
- The reward may be stochastic.
- Reward is delayed (i.e. finding food in a maze)

Complications

- You may have no clue (model) of how rewards are being paid off.
- The world may change while you try to learn it
- How much time do you need to explore uncharted territory before you exploit what you have learned?

The Task

- To learn an optimal *policy* that maps states of the world to actions of the agent.

I.e., if this patch of room is dirty, I clean it. If my battery is empty, I recharge it.

$$\pi : \mathcal{S} \rightarrow \mathcal{A}$$

- What is it that the agent tries to optimize?
Answer: the total future discounted reward:

The Task

- What is it that the agent tries to optimize?

Answer: the **total future discounted reward**:

$$\begin{aligned} V^\pi(s_t) &= \underbrace{r_t}_{\text{immediate}} + \underbrace{\gamma r_{t+1}}_{\text{discounted}} + \underbrace{\gamma^2 r_{t+2}}_{\text{discounted}} + \dots \\ &= \sum_{i=0}^{\infty} \gamma^i r_{t+i} \quad \underbrace{0 \leq \gamma < 1}_{\text{discount factor}} \end{aligned}$$

Note: immediate reward is worth more than future reward.

What would happen to mouse in a maze with gamma = 0 ?

Value Function

- Let's say we have access to the optimal value function that computes the total future discounted reward $V^*(s)$


- What would be the optimal policy $\pi^*(s)$?

- Answer: we choose the action that maximizes:

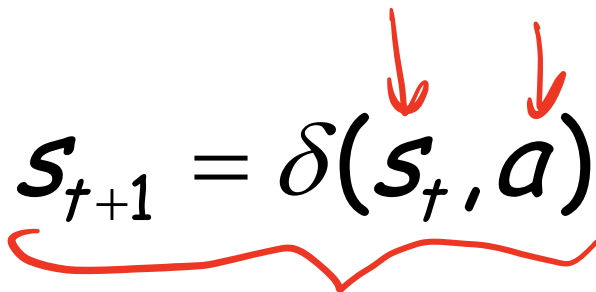
$$\pi^*(s) = \arg \max_a \left[\underbrace{r(s, a)}_{\substack{\downarrow \downarrow \\ s_{t+1} = \delta(s_t, a)}} + \gamma V^*(\delta(s, a)) \right]$$

Value Function

- We assume that we know what the reward will be if we perform action “a” in state “s”:


$$r(s, a)$$


- We also assume we know what the next state of the world will be if we perform action “a” in state “s”:

$$s_{t+1} = \delta(s_t, a)$$


Example I

- Consider some complicated graph, and we would like to find the shortest path from a node S_i to a goal node G .
- Traversing an edge will cost you “length edge” dollars.




Example I

- The value function encodes the total remaining distance to the goal node from any node s , i.e.

$$V(s) = \text{"1 / distance"} \text{ to goal from } s.$$

- If you know $V(s)$, the problem is trivial. You simply choose the node that has highest $V(s)$.




Example II


- A simple deterministic world.
- Each grid square represents a distinct state, each arrow a distinct action.
- The immediate reward function, $r(s, a)$ gives reward 100 for actions entering the goal state G , and zero otherwise. Values of $V^*(s)$ follow from $r(s, a)$, and the discount factor $\gamma = 0.9$.
- An optimal policy is also shown.

Example II


Find your way to the goal.



$r(s, a)$ (immediate reward) values



$V^*(s)$ values



One optimal policy

$$\begin{aligned}
 V^\pi(s_t) &= r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots \\
 &= \sum_{i=0}^{\infty} \gamma^i r_{t+i} \quad 0 \leq \gamma < 1
 \end{aligned}$$

Q-Function

- One approach to RL is then to try to estimate $V^*(s)$.

Bellman Equation:

$$V^*(s) \leftarrow \max_a [r(s,a) + \gamma V^*(\delta(s,a))]$$

- However, this approach requires you to know $r(s,a)$ and $\delta(s,a)$.
- This is unrealistic in many real problems. What is the reward if a robot is exploring mars and decides to take a right turn?

Q-Function

- Fortunately we can circumvent this problem by exploring and experiencing how the world reacts to our actions. We need to *learn* r & δ .

Q-Function

- We want a function that directly learns good state-action pairs, i.e. what action should I take in this state. We call this $Q(s,a)$.

Q-Function

- Let us define the evaluation function $Q(s, a)$ so that its value is the maximum discounted cumulative reward that can be achieved starting from state s and applying action a as the first action.
- In other words, the value of Q is the reward received immediately upon executing action a from state s , plus the value (discounted by γ) of following the optimal policy thereafter.

$$Q(s, a) \equiv r(s, a) + \gamma V^*(\delta(s, a))$$

Q-Function

- Why is this rewrite important? Because it shows that if the agent learns the Q function instead of the V^* function, it will be able to select optimal actions even when it has no knowledge of the functions r and δ .

Q-Function

- It need only consider each available action a in its current state s and choose the action that maximizes $Q(s, a)$.


$$\pi^*(s) = \operatorname{argmax}_a Q(s, a)$$

$$V^*(s) = \max_a Q(s, a)$$


Example

- To illustrate, the figure in the next slide shows the Q values for every state and action in the simple grid world.
- The Q value for each state-action transition equals the r value for this transition plus the V^* value for the resulting state discounted by γ .
- The optimal policy shown in the figure corresponds to selecting actions with maximal Q values.


Example II



$V^*(s)$ values



$Q(s, a)$ values



One optimal policy

Check that

$$\pi^*(s) = \arg \max_a Q(s, a)$$

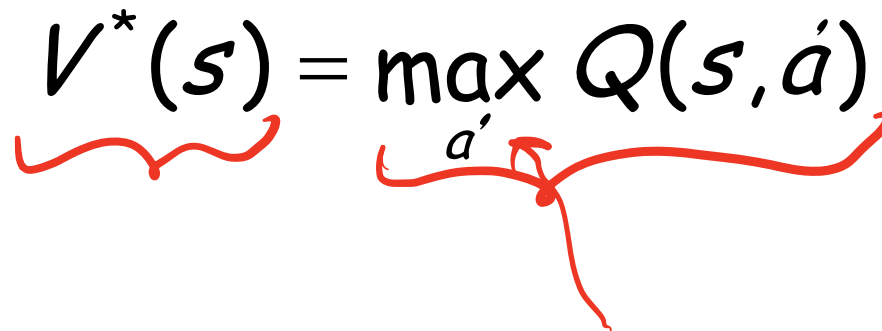
$$V^*(s) = \max_a Q(s, a)$$

Q-Learning

- Learning the Q function corresponds to learning the optimal policy. How can Q be learned?
- The key problem is finding a reliable way to estimate training values for Q , given only a sequence of immediate rewards r spread out over time.

Q-Learning

- This can be accomplished through iterative approximation.
- To see how, notice the close relationship between Q and V^* ,

$$V^*(s) = \max_{a'} Q(s, a')$$


Q-Learning

which allows rewriting the Q function as:

$$\begin{aligned} Q(s, a) &\equiv \underbrace{r(s, a)} + \underbrace{\gamma V^*(\delta(s, a))} \\ &= r(s, a) + \gamma \max_{\underbrace{a'}} \underbrace{Q(\delta(s, a), a')} \end{aligned}$$

This still depends on $r(s, a)$ and $\delta(s, a)$;
however,

Q-Learning

$$\begin{aligned} Q(s, a) &\equiv r(s, a) + \gamma V^*(\delta(s, a)) \\ &= r(s, a) + \gamma \max_{a'} Q(\delta(s, a), a') \end{aligned}$$

this recursive definition of Q provides the basis for algorithms that iteratively approximate Q .

\hat{Q} refers to the learner's estimate, or hypothesis, of the actual Q function.

Q-Learning

$$\begin{aligned} Q(s, a) &\equiv r(s, a) + \gamma V^*(\delta(s, a)) \\ &= r(s, a) + \gamma \max_{a'} Q(\delta(s, a), a') \end{aligned}$$

- Imagine the robot is exploring its environment, trying new actions as it goes.
- At every step it receives some reward “ r ”, and it observes the environment change into a new state s' for action a .
- How can we use these observations, (s, a, s', r) to learn a model?



Q-Learning


- The learner represents its hypothesis \hat{Q} by a large table with a separate entry for each state-action pair.
- The table entry for the pair (s, a) stores the value for $\hat{Q}(s, a)$, learner's current hypothesis about the actual but unknown value $Q(s, a)$.
- The table can be initially filled with random values (though it is easier to understand the algorithm if one assumes initial values of zero).

Q-Learning

- The agent repeatedly observes its current state s , chooses some action a , executes this action, then observes the resulting reward $r = r(s, a)$ and the new state $s' = \delta(s, a)$.

Q-Learning

- It then updates the table entry for $\hat{Q}(s,a)$ following each such transition, according to the rule:

$$\hat{Q}(s,a) \leftarrow r + \gamma \max_{a'} \hat{Q}(s',a') \quad s' = s_{t+1}$$


Q-Learning

$$\hat{Q}(s, a) \leftarrow r + \gamma \max_{a'} \hat{Q}(s', a') \quad s' = s_{t+1}$$

- This equation continually estimates Q at state s consistent with an estimate of Q at state s' , one step in the future: temporal difference (TD) learning.
- Note that s' is closer to goal, and hence more “reliable”, but still an estimate itself.

Q-Learning Summary

Q learning algorithm

For each s, a initialize the table entry $\hat{Q}(s, a)$ to zero.

Observe the current state s

Do forever:

- Select an action a and execute it
- Receive immediate reward r
- Observe the new state s'
- Update the table entry for $\hat{Q}(s, a)$ as follows:

$$\hat{Q}(s, a) \leftarrow r + \gamma \max_{a'} \hat{Q}(s', a')$$


- $s \leftarrow s'$

Q-Learning

$$\hat{Q}(s, a) \leftarrow r + \gamma \max_{a'} \hat{Q}(s', a') \\ s' = s_{t+1}$$

- We do an update after each state-action pair. I.e., we are learning online!
- We are learning useful things about explored state-action pairs. These are typically most useful because they are likely to be encountered again.
- Under suitable conditions, these updates can actually be proved to converge to the real answer.

Example: Q-Learning



$$\hat{Q}(s_1, a_{right}) \leftarrow r + \gamma \max_{a'} \hat{Q}(s_2, a')$$

$$\leftarrow 0 + 0.9 \max\{63, 81, 100\}$$

$$\leftarrow 90$$

Q-learning propagates Q-estimates 1-step backwards

Exploration / Exploitation

- The algorithm does not specify how actions are chosen by the agent.
- One strategy would be in state s to select the action a that maximizes $\hat{Q}(s, a)$, thereby *exploiting* its current approximation \hat{Q} .

Exploration / Exploitation

- Using this strategy the agent runs the risk that it will overcommit to actions that are found during early training to have high Q values, while failing to explore other actions that have even higher values.
- It is common in Q learning to use a probabilistic approach to selecting actions.

Simulated Annealing Exploration / Exploitation

- Actions with higher Q values are assigned higher probabilities, but every action is assigned a nonzero probability. One way to assign such probabilities is

$$P(a_i | s) = \frac{e^{\hat{Q}(s, a_i)/T}}{\sum_j e^{\hat{Q}(s, a_j)/T}}$$

Handwritten notes:

- $T \rightarrow 1$: The higher the \hat{Q} , the larger $P(a_i | s)$
- $T \rightarrow \infty$: Uniform \hat{Q} doesn't matter

- where $P(a_i | s)$ is the probability of selecting action a_i , given that the agent is in state s , and where $T > 0$ is a constant that determines how strongly the selection favors actions with high Q values.

Exploration / Exploitation

- Hence it is good to try new things so now and then, e.g. If T large lots of **exploring**, if T small, **exploit** current policy. One can decrease T over time to first explore, and then converge and exploit.
- For example $T = c/k + d$ where k is iteration of the algorithm

$$P(a_i | s) = \frac{e^{\hat{Q}(s, a_i)/T}}{\sum_j e^{\hat{Q}(s, a_j)/T}}$$

- Decreasing T over time is sometimes called **simulated annealing**, which is inspired by annealing process in metals. T is sometimes called the **Temperature**.

Improvements

- One can trade-off memory and computation by caching (s, s', r) for observed transitions. After a while, as $Q(s', a')$ has changed, you can “replay” the update:

$$\hat{Q}(s, a) \leftarrow r + \gamma \max_{a'} \hat{Q}(s', a')$$

- One can actively search for state-action pairs for which $Q(s, a)$ is expected to change a lot (prioritized sweeping).

Temporal Difference Learning

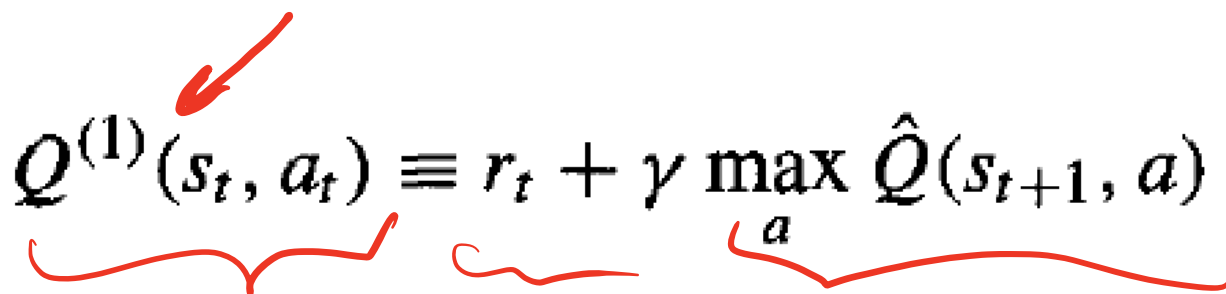
- Q learning algorithm learns by iteratively reducing the discrepancy between Q value estimates for adjacent state
- Q learning is a special case of *temporal difference* algorithms that learn by reducing discrepancies between estimates made by the agent at different times.

Temporal Difference Learning

- The training rule we studied reduces the difference between the estimated Q values of a state and its immediate successor
- However, we could design an algorithm that reduces discrepancies between this state and more distant descendants or ancestors.

Temporal Difference Learning

- Recall that our Q learning training rule calculates a training value for $\hat{Q}(s_t, a_t)$ in terms of the values for $\hat{Q}(s_{t+1}, a_{t+1})$ where s_{t+1} is the result of applying action a_t to the state s_t .
- Let $Q^{(1)}(s_t, a_t)$ denote the training value calculated by this one-step lookahead:

$$Q^{(1)}(s_t, a_t) \equiv r_t + \gamma \max_a \hat{Q}(s_{t+1}, a)$$


Temporal Difference Learning

- One alternative way to compute a training value for $Q(s_t, a_t)$ is to base it on the observed rewards for two steps


$$Q^{(2)}(s_t, a_t) \equiv r_t + \gamma r_{t+1} + \gamma^2 \max_a \hat{Q}(s_{t+2}, a)$$

- or, in general, for n steps

$$Q^{(n)}(s_t, a_t) \equiv r_t + \gamma r_{t+1} + \cdots + \gamma^{(n-1)} r_{t+n-1} + \gamma^n \max_a \hat{Q}(s_{t+n}, a)$$

Temporal Difference Learning

- A general method for blending these alternative training estimates, called TD(λ).
- The idea is to use a constant $0 \leq \lambda \leq 1$ to combine the estimates obtained from various lookahead distances in the following fashion

$$Q^\lambda(s_t, a_t) \equiv (1 - \lambda) \left[\underbrace{Q^{(1)}(s_t, a_t)} + \underbrace{\lambda Q^{(2)}(s_t, a_t)} + \underbrace{\lambda^2 Q^{(3)}(s_t, a_t)} + \dots \right]$$


Temporal Difference Learning

- An equivalent recursive definition for Q^λ is

$$Q^\lambda(s_t, a_t) = \underbrace{r_t}_{\text{reward}} + \gamma \left[\underbrace{(1 - \lambda)}_{\text{discount}} \underbrace{\max_a \hat{Q}(s_t, a_t)}_{\text{max over actions}} + \underbrace{\lambda Q^\lambda(s_{t+1}, a_{t+1})}_{\text{bootstrapped estimate}} \right]$$

$\lambda = 0 \Rightarrow$ back to Q -learning

Temporal Difference Learning

$$Q^\lambda(s_t, a_t) = r_t + \gamma [(1 - \lambda) \max_a \hat{Q}(s_t, a_t) + \lambda Q^\lambda(s_{t+1}, a_{t+1})]$$

- If $\lambda = 0$ we have our original training estimate $Q^{(1)}$, which considers only one-step discrepancies in the Q estimates.
- As λ is increased, the algorithm places increasing emphasis on discrepancies based on more distant lookaheads.

Temporal Difference Learning

$$Q^\lambda(s_t, a_t) = r_t + \gamma [(1 - \lambda) \max_a \hat{Q}(s_t, a_t) + \lambda Q^\lambda(s_{t+1}, a_{t+1})]$$

- At the extreme value $\lambda = 1$, only the observed r_{t+i} values are considered, with no contribution from the current Q estimate.
- The motivation for the TD(λ) method is that in some settings training will be more efficient if more distant lookaheads are considered.

Extensions

$$E_{s'} [\max_{a'} Q(s', a')]$$

- To deal with stochastic environments, we need to maximize **expected** future discounted reward:

$$Q(s, a) = E[r(s, a)] + \gamma \sum_{s'} P(s' | s, a) \max_{a'} Q(s', a')$$

- Often the state space is **too large** to deal with all states and adopt a table-lookup approach. In this case we need to learn a function:

$$Q(s, a) \approx f_{\theta}(s, a)$$

Extensions

$$Q(s, a) \approx f_{\theta}(s, a)$$

- Neural network with back-propagation have been quite successful.

- For instance, TD-Gammon is a back-gammon program that plays at expert level.
- state-space very large, trained by playing against itself, uses NN to approximate value function, uses $TD(\lambda)$ for learning.

More on Function Approximation

- For instance: linear function:

$$Q(s, a) \approx f_{\theta}(s, a) = \sum_{k=1}^K \theta_k^a \Phi_k(s)$$

The features Φ are fixed measurements of the state (e.g. # stones on the board).

We only learn the parameters theta.

- Update rule: (start in state s , take action a , observe reward r and end up in state s')

$$\theta_k^a \leftarrow \theta_k^a + \alpha \left(r + \gamma \max_{a'} \hat{Q}(s', a') - \hat{Q}(s, a) \right) \Phi_k(s)$$

change in Q

Conclusion

- Reinforcement learning addresses a very broad and relevant question:
How can we learn to survive in our environment?
- We have looked at Q-learning, which simply learns from experience.
No model of the world is needed.

Conclusion

- We made simplifying assumptions: e.g. state of the world only depends on last state and action. This is the *Markov* assumption. The model is called a *Markov Decision Process (MDP)*.

Conclusion

- We assumed deterministic dynamics, reward function, but the world really is stochastic.
- There are many extensions to speed up learning.
- There have been many successful real world applications.